

# AML\_HW4\_Questions\_th3061

November 17, 2023

## 0.0.1 Instructions

- 1) Please push the .ipynb and .pdf to Github Classroom prior to the deadline, .py file is optional (not needed).
- 2) Please include your Name and UNI below.

**0.0.2 Name: Ta-Wei Huang**

**0.0.3 UNI: th3061**

##Setup

```
[ ]: import numpy as np
import matplotlib.pyplot as plt
import pprint
pp = pprint.PrettyPrinter(indent=4)
import warnings
warnings.filterwarnings("ignore")
```

## 1 Part 1: Neural Network from the scratch

For this part, you are not allowed to use any library other than numpy.

In this part, you will implement the forward pass and backward pass (i.e. the derivatives of each parameter wrt to the loss) with the network image uploaded.

The weight matrix for the hidden layer is  $W1$  and has bias  $b1$ .

The weight matrix for the output layer is  $W2$  and has bias  $b2$ .

Activation function is sigmoid for both hidden and output layer

Loss function is the MSE loss

Refer to the below dictionary for dimensions for each matrix

```
[ ]: np.random.seed(0) # don't change this
weights = {
    'W1': np.random.randn(3, 2),
    'b1': np.zeros(3),
    'W2': np.random.randn(3),
    'b2': 0,
```

```

}
X = np.random.rand(1000,2)
Y = np.random.randint(low=0, high=2, size=(1000,))

```

```
[ ]: weights
```

```
[ ]: {'W1': array([[ 1.76405235,  0.40015721],
                  [ 0.97873798,  2.2408932 ],
                  [ 1.86755799, -0.97727788]]),
      'b1': array([0., 0., 0.]),
      'W2': array([ 0.95008842, -0.15135721, -0.10321885]),
      'b2': 0}
```

```
[ ]: #Sigmoid Function
def sigmoid(z):
    return 1/(1 + np.exp(-z))

```

```
[ ]: #Implement the forward pass
def forward_propagation(X, weights):
    # Z1 -> output of the hidden layer before applying activation
    # H -> output of the hidden layer after applying activation
    # Z2 -> output of the final layer before applying activation
    # Y -> output of the final layer after applying activation

    Z1 = np.dot(X, weights['W1'].T) + weights['b1']
    H = sigmoid(Z1)

    Z2 = np.dot(H, weights['W2'].T) + weights['b2']
    Y = sigmoid(Z2)

    return Y, Z2, H, Z1

```

```
[ ]: # Implement the backward pass
# Y_T are the ground truth labels
def back_propagation(X, Y_T, weights):
    N_points = X.shape[0]

    # forward propagation
    Y, Z2, H, Z1 = forward_propagation(X, weights)
    L = (1/(2*N_points)) * np.sum(np.square(Y - Y_T))

    # back propagation
    dLdY = 1/N_points * (Y - Y_T)
    dLdZ2 = np.multiply(dLdY, (sigmoid(Z2)*(1-sigmoid(Z2))))
    dLdW2 = np.dot(H.T, dLdZ2)

    ones = np.ones((1000))

```

```

dLdb2 = np.dot(ones.T, dLdZ2)
dLdH = np.dot(dLdZ2.reshape(-1,1), weights['W2'].reshape(-1,1).T)
# dLdZ1 =
dLdZ1 = np.multiply(dLdH, (sigmoid(Z1)*(1-sigmoid(Z1))))
# dLdW1 =
dLdW1 = np.dot(X.T, dLdZ1)
# dLdb1 =
dLdb1 = np.dot(ones.T, dLdZ1)

gradients = {
    'W1': dLdW1,
    'b1': dLdb1,
    'W2': dLdW2,
    'b2': dLdb2,
}

return gradients, L

```

```

[ ]: gradients, L = back_propagation(X, Y, weights)
print(L)

```

0.1332476222330792

```

[ ]: pp.pprint(gradients)

```

```

{   'W1': array([[ 0.00244596, -0.00030765, -0.00034768],
                 [ 0.00262019, -0.00024188, -0.000372   ]]),
    'W2': array([0.02216011, 0.02433097, 0.01797174]),
    'b1': array([ 0.00492577, -0.00058023, -0.00065977]),
    'b2': 0.029249230265318685}

```

Your answers should be close to  $L = 0.133$  and `'b1': array([ 0.00492, -0.000581, -0.00066])`. You will be graded based on your implementation and outputs for  $L$ ,  $W1$ ,  $W2$ ,  $b1$ , and  $b2$

## 2 Part 2: Neural network to classify images: CIFAR-10

CIFAR-10 is a dataset of 60,000 color images (32 by 32 resolution) across 10 classes (airplane, automobile, bird, cat, deer, dog, frog, horse, ship, truck). The train/test split is 50k/10k.

```

[ ]: from tensorflow.keras.datasets import cifar10
(x_dev, y_dev), (x_test, y_test) = cifar10.load_data()

LABELS = 
↪ ['airplane', 'automobile', 'bird', 'cat', 'deer', 'dog', 'frog', 'horse', 'ship', 'truck']

```

Downloading data from <https://www.cs.toronto.edu/~kriz/cifar-10-python.tar.gz>  
170498071/170498071 [=====] - 5s 0us/step

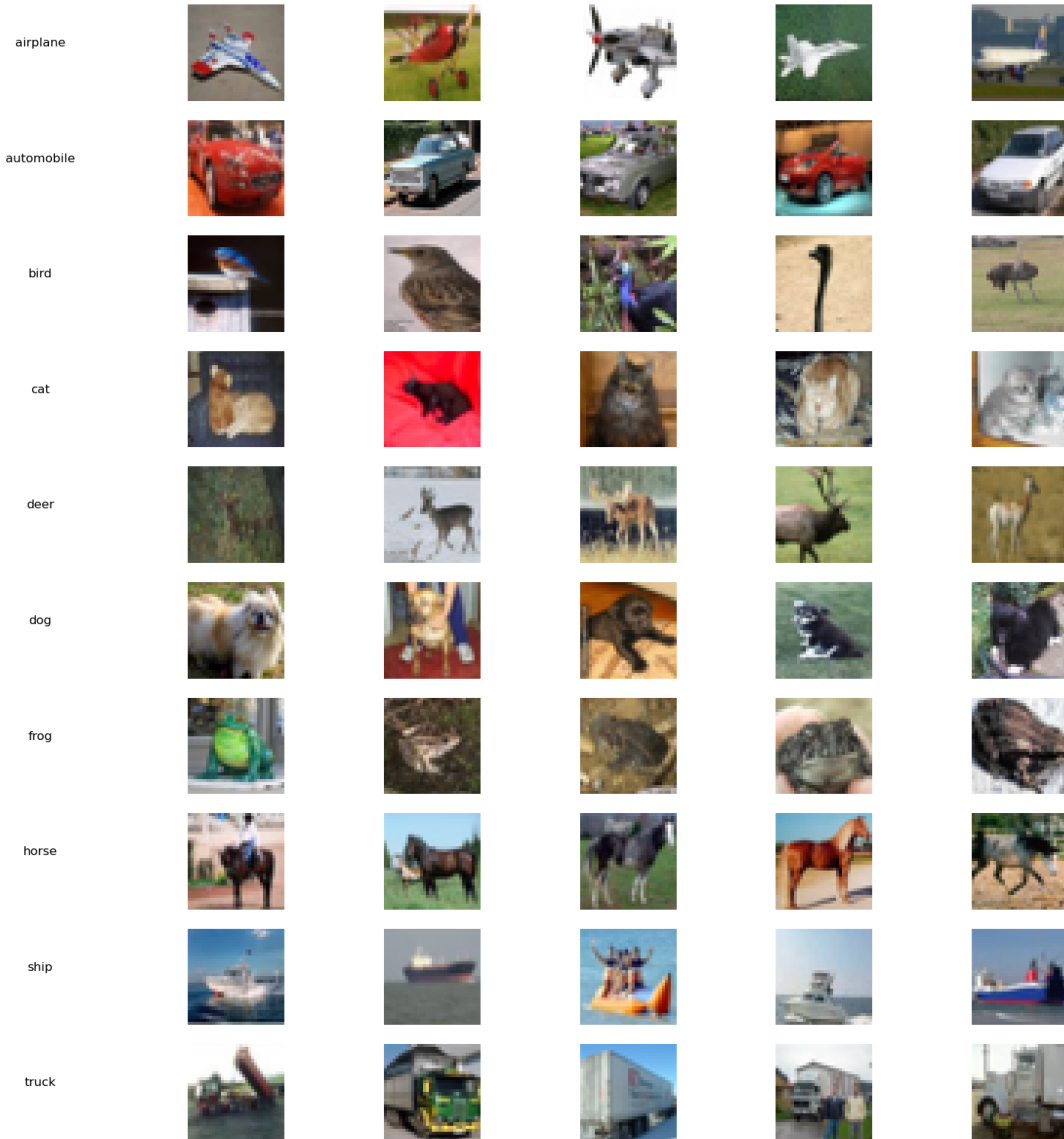
### 2.0.1 2.1 Plot 50 samples from each class/label from train set on a 10\*5 subplot

```
[ ]: #Your code here
def plot_samples(x, y, labels, num_samples=5, num_classes=10):
    fig, axes = plt.subplots(num_classes, num_samples+1, figsize=(20, 20))
    #fig.subplots_adjust(hspace=0.5, wspace=0.1)

    for i, label in enumerate(labels):
        idx = np.where(y == i)[0][:num_samples+1 ]

        for j, img_idx in enumerate(idx):
            ax = axes[i, j]
            if j == 0:
                ax.set_title(label, size=12, y=0.5)
                ax.axis('off')
            else:
                ax.imshow(x[img_idx])
                ax.axis('off')

plot_samples(x_dev, y_dev, LABELS)
plt.show()
```



## 2.0.2 2.2 Preparing the dataset for NN

- 1) Print the shapes - , , ,
- 2) Flatten the images into one-dimensional vectors and again print the shapes of ,
- 3) Standardize the development and test sets.
- 4) Train-test split your development set into train and validation sets (8:2 ratio).

```
[ ]: # 1
print("x_dev shape:", x_dev.shape) # number, height, width, color
print("y_dev shape:", y_dev.shape)
print("x_test shape:", x_test.shape)
```

```

print("y_test shape:", y_test.shape)
# 2
x_dev_flat = x_dev.reshape(x_dev.shape[0], -1) # 50000*1
x_test_flat = x_test.reshape(x_test.shape[0], -1)
print("x_dev_flat shape:", x_dev_flat.shape)
print("x_test_flat shape:", x_test_flat.shape)
# 3
x_dev_standardized = x_dev_flat / 255.0
x_test_standardized = x_test_flat / 255.0
# 4
from sklearn.model_selection import train_test_split
x_train, x_val, y_train, y_val = train_test_split(x_dev_standardized, y_dev,
    ↪test_size=0.2, random_state=42)

print("x_train shape:", x_train.shape)
print("x_val shape:", x_val.shape)
print("y_train shape:", y_train.shape)
print("y_val shape:", y_val.shape)

```

```

x_dev shape: (50000, 32, 32, 3)
y_dev shape: (50000, 1)
x_test shape: (10000, 32, 32, 3)
y_test shape: (10000, 1)
x_dev_flat shape: (50000, 3072)
x_test_flat shape: (10000, 3072)
x_train shape: (40000, 3072)
x_val shape: (10000, 3072)
y_train shape: (40000, 1)
y_val shape: (10000, 1)

```

### 2.0.3 2.3 Build the feed forward network with the below specifications

First layer size = 128

hidden layer size = 64

last layer size = Figure this out from the data!

```

[ ]: import tensorflow as tf
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense, Flatten

model = Sequential([
    Dense(128, activation='relu', input_shape=(3072,)), # First dense layer
    ↪with 128 units
    Dense(64, activation='relu'), # Second (hidden)
    ↪dense layer with 64 units
])

```

```
Dense(10, activation='softmax') # Output layer with 10 units for 10 classes
])
```

#### 2.0.4 2.4 Print out the model summary. Mention the number of parameters for each layer.

```
[ ]: model.summary()
```

Model: "sequential"

Layer (type)	Output Shape	Param #
dense (Dense)	(None, 128)	393344
dense_1 (Dense)	(None, 64)	8256
dense_2 (Dense)	(None, 10)	650

=====  
 Total params: 402250 (1.53 MB)  
 Trainable params: 402250 (1.53 MB)  
 Non-trainable params: 0 (0.00 Byte)

Layer (type)	Output Shape	Param #
dense (Dense)	(None, 128)	393344
dense_1 (Dense)	(None, 64)	8256
dense_2 (Dense)	(None, 10)	650

=====  
 Total params: 402250 (1.53 MB)  
 Trainable params: 402250 (1.53 MB)  
 Non-trainable params: 0 (0.00 Byte)

#### 2.0.5 2.5 Do you think the number of parameters is dependent on the image height and width?

Yes, the number of parameters in NN is dependent on the image height and width. For example, in the initial layer, the number of parameters is `input_size * number_of_neuron + number_of_neuron(bias)`, as for each neuron, there is one weight per input unit.

Printing out your model's output on first train sample. This will confirm if your dimensions are correctly set up. The sum of this output equal to 1 upto two decimal

places?

```
[ ]: #modify name of X_train based on your requirement
```

```
model.compile()
output = model.predict(x_train[0].reshape(1,-1))

print("Output: {:.2f}".format(sum(output[0])))
```

```
1/1 [=====] - 0s 24ms/step
```

```
Output: 1.00
```

```
1/1 [=====] - 0s 24ms/step
```

```
Output: 1.00
```

**2.0.6 2.6 Using the right metric and the right loss function, with Adam as the optimizer, train your model for 20 epochs.**

```
[ ]: model.compile(optimizer='adam',
                  loss='sparse_categorical_crossentropy',
                  metrics=['accuracy'])
history = model.fit(x_train, y_train, epochs=20, validation_split=0.2)
```

```
Epoch 1/20
```

```
1000/1000 [=====] - 1s 1ms/step - loss: 1.2801 -
```

```
accuracy: 0.5429 - val_loss: 1.3402 - val_accuracy: 0.5232
```

```
Epoch 2/20
```

```
1000/1000 [=====] - 1s 895us/step - loss: 1.2733 -
```

```
accuracy: 0.5417 - val_loss: 1.3265 - val_accuracy: 0.5321
```

```
Epoch 3/20
```

```
1000/1000 [=====] - 1s 935us/step - loss: 1.2667 -
```

```
accuracy: 0.5471 - val_loss: 1.3578 - val_accuracy: 0.5217
```

```
Epoch 4/20
```

```
1000/1000 [=====] - 1s 894us/step - loss: 1.2632 -
```

```
accuracy: 0.5493 - val_loss: 1.3722 - val_accuracy: 0.5176
```

```
Epoch 5/20
```

```
1000/1000 [=====] - 1s 894us/step - loss: 1.2554 -
```

```
accuracy: 0.5513 - val_loss: 1.4578 - val_accuracy: 0.4846
```

```
Epoch 6/20
```

```
1000/1000 [=====] - 1s 923us/step - loss: 1.2599 -
```

```
accuracy: 0.5469 - val_loss: 1.3771 - val_accuracy: 0.5141
```

```
Epoch 7/20
```

```
1000/1000 [=====] - 1s 955us/step - loss: 1.2479 -
```

```
accuracy: 0.5532 - val_loss: 1.3708 - val_accuracy: 0.5214
```

```
Epoch 8/20
```

```
1000/1000 [=====] - 1s 899us/step - loss: 1.2451 -
```

```
accuracy: 0.5542 - val_loss: 1.3787 - val_accuracy: 0.5131
```

```
Epoch 9/20
```

```
1000/1000 [=====] - 1s 917us/step - loss: 1.2401 -
```

```
accuracy: 0.5536 - val_loss: 1.3948 - val_accuracy: 0.5176
```



```

Epoch 10/20
1000/1000 [=====] - 1s 906us/step - loss: 1.2353 -
accuracy: 0.5562 - val_loss: 1.4000 - val_accuracy: 0.5121
Epoch 11/20
1000/1000 [=====] - 1s 952us/step - loss: 1.2314 -
accuracy: 0.5574 - val_loss: 1.4313 - val_accuracy: 0.4985
Epoch 12/20
1000/1000 [=====] - 1s 922us/step - loss: 1.2270 -
accuracy: 0.5610 - val_loss: 1.3922 - val_accuracy: 0.5113
Epoch 13/20
1000/1000 [=====] - 1s 899us/step - loss: 1.2284 -
accuracy: 0.5595 - val_loss: 1.4279 - val_accuracy: 0.5024
Epoch 14/20
1000/1000 [=====] - 1s 950us/step - loss: 1.2252 -
accuracy: 0.5578 - val_loss: 1.4062 - val_accuracy: 0.5107
Epoch 15/20
1000/1000 [=====] - 1s 971us/step - loss: 1.2193 -
accuracy: 0.5623 - val_loss: 1.4261 - val_accuracy: 0.5023
Epoch 16/20
1000/1000 [=====] - 1s 906us/step - loss: 1.2140 -
accuracy: 0.5652 - val_loss: 1.4303 - val_accuracy: 0.5107
Epoch 17/20
1000/1000 [=====] - 1s 978us/step - loss: 1.2176 -
accuracy: 0.5628 - val_loss: 1.4091 - val_accuracy: 0.5109
Epoch 18/20
1000/1000 [=====] - 1s 905us/step - loss: 1.2125 -
accuracy: 0.5661 - val_loss: 1.4694 - val_accuracy: 0.4935
Epoch 19/20
1000/1000 [=====] - 1s 903us/step - loss: 1.2129 -
accuracy: 0.5643 - val_loss: 1.4343 - val_accuracy: 0.5044
Epoch 20/20
1000/1000 [=====] - 1s 1ms/step - loss: 1.2094 -
accuracy: 0.5676 - val_loss: 1.4229 - val_accuracy: 0.5092

```

## 2.0.7 2.7 Plot the training curves described below

```

[ ]: import matplotlib.pyplot as plt

train_loss = history.history['loss']
epochs = range(1, len(train_loss) + 1)

plt.figure()
plt.plot(epochs, train_loss, 'r', label='Training loss')
plt.title('Training Loss')
plt.xlabel('Epochs')
plt.ylabel('Loss')
plt.legend()

```

```
plt.show()
```



## 2.0.8 2.7.1 Display the train vs validation loss over each epoch

```
[ ]: import matplotlib.pyplot as plt

train_loss = history.history['loss']
val_loss = history.history['val_loss']
epochs = range(1, len(train_loss) + 1)


plt.figure()
plt.plot(epochs, train_loss, 'r', label='Training loss')
plt.plot(epochs, val_loss, 'b', label='Validation loss')
plt.title('Training and Validation Loss')
plt.xlabel('Epochs')
plt.ylabel('Loss')
plt.legend()
plt.show()
```



### 2.0.9 2.7.2 Display the train vs validation accuracy over each epoch

```
[ ]: train_acc = history.history['accuracy']
      val_acc = history.history['val_accuracy']
      epochs = range(1, len(train_acc) + 1)

      plt.figure()
      plt.plot(epochs, train_acc, 'r', label='Training acc')
      plt.plot(epochs, val_acc, 'b', label='Validation acc')
      plt.title('Training and Validation Accuracy')
      plt.xlabel('Epochs')
      plt.ylabel('Accuracy')
      plt.legend()
      plt.show()
```



#### 2.0.10 2.8 Finally, report the metric chosen on test set

```
[ ]: test_loss, test_acc = model.evaluate(x_test_standardized, y_test, verbose=2)
      print('Test accuracy:', test_acc)
```

```
313/313 - 0s - loss: 1.4972 - accuracy: 0.4846 - 172ms/epoch - 548us/step
Test accuracy: 0.4846000075340271
Test accuracy: 0.4846000075340271
```

#### 2.0.11 2.9 Plot the first 50 samples of test dataset on a 10\*5 subplot and this time label the images with both the ground truth (GT) and predicted class (P). (Make sure you predict the class with the improved model)

```
[ ]: predictions = model.predict(x_test_flat[:50])
      #print(predictions)
      predicted_classes = np.argmax(predictions, axis=1)
      #print(predicted_classes)
      fig, axes = plt.subplots(10, 5, figsize=(15, 30))
      fig.subplots_adjust(hspace=0.5, wspace=0.1)
      for i, ax in enumerate(axes.flat):
```

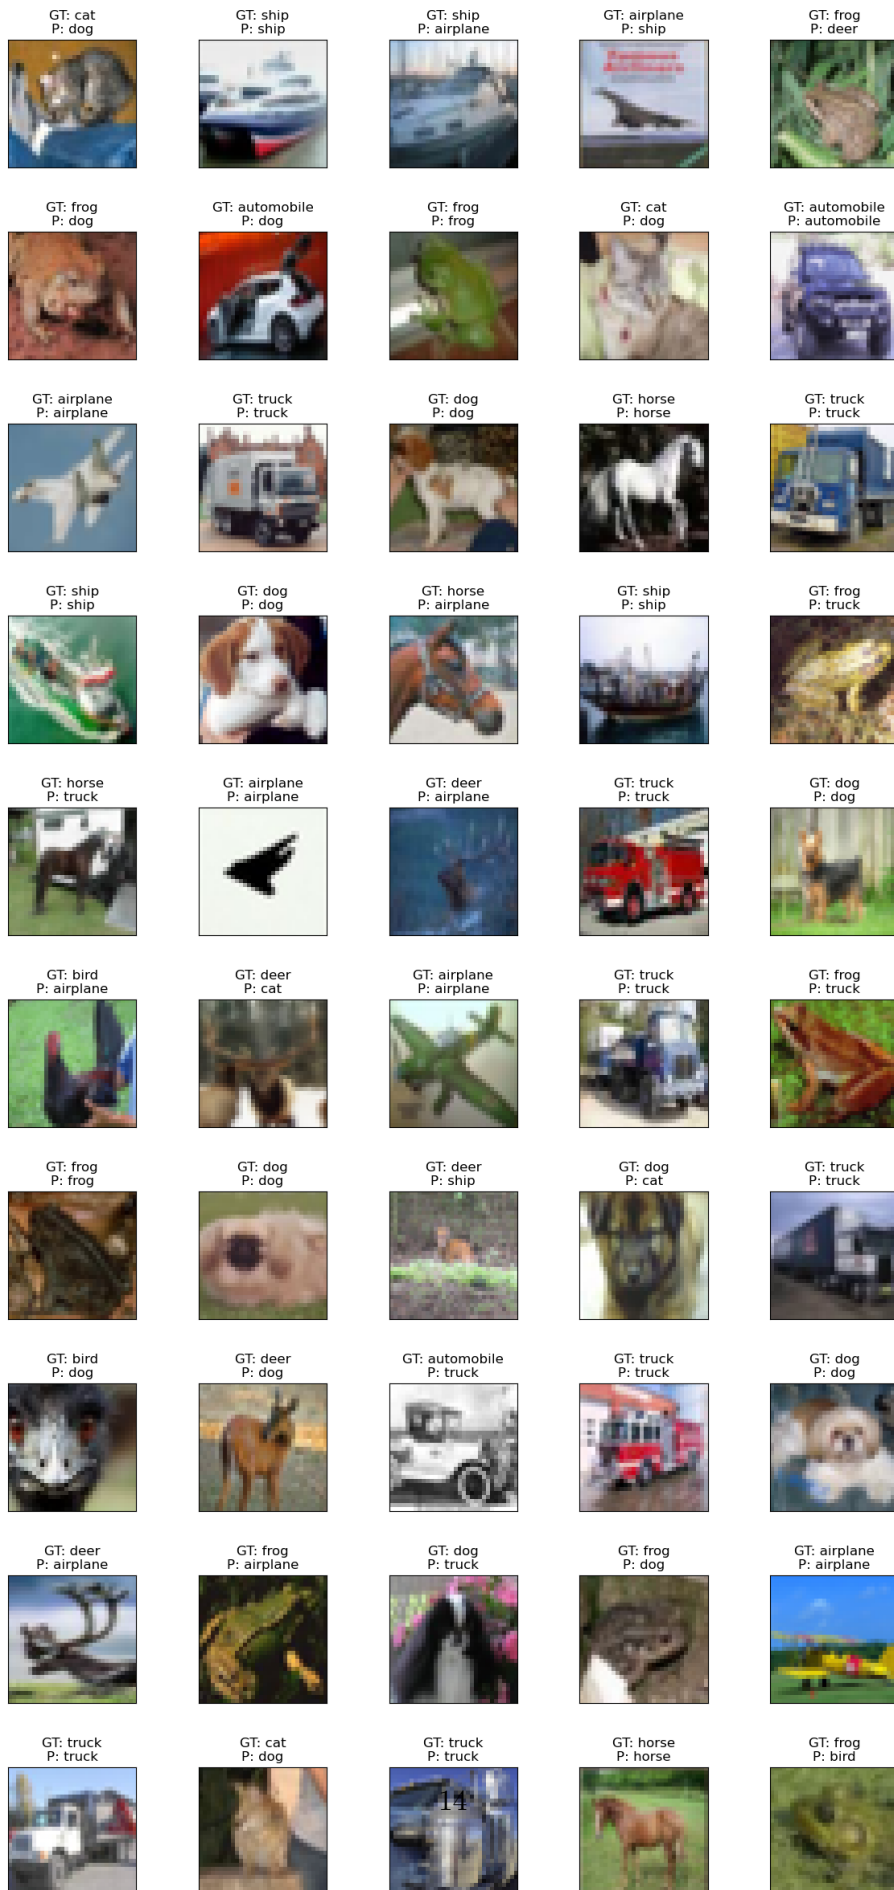
```
# Plot image
ax.imshow(x_test[i])

# Add label with ground truth and prediction
true_label = LABELS[y_test[i][0]]
predicted_label = LABELS[predicted_classes[i]]
ax.set_title(f"GT: {true_label}\nP: {predicted_label}")

# Remove axis ticks
ax.set_xticks([])
ax.set_yticks([])

plt.show()
```

2/2 [=====] - 0s 1ms/step



[ ]: